

Computer Vision Knife Detection

April 20, 2018

0.1 Computer Vision Knife Detection

```
In [2]: import matplotlib.pyplot as plt
        from urllib.request import urlretrieve
        from os.path import isfile, isdir, getsize
        from os import mkdir, makedirs, remove, listdir
        from tqdm import tqdm

        import zipfile
        import tarfile
        import imageio
        import glob
        import shutil
        import pickle
        import numpy as np

        from keras.models import Sequential, Model
        from keras import optimizers
        from keras.layers import Dense, Activation, Conv2D, MaxPool2D, Flatten, BatchNormalization
        from keras.preprocessing.image import ImageDataGenerator
```

0.1.1 Webcam Initialization

```
In [3]: import cv2

        camera = cv2.VideoCapture(0)

        def get_image():
            re, im = camera.read()
            return im
```

0.1.2 Dataset Initialization

```
In [4]: datasetfolderpath = "/Users/kevin/desktop/Datas/knives"
        files = glob.glob(datasetfolderpath + "/data/**/*.bmp")

        labels = np.array([0]*3560 + [1]*9341)
```

```

In [5]: len_data = len(files)
        train_examples = 10000
        test_examples = len_data - train_examples

        permutation = np.random.permutation(len_data)
        train_set = [files[i] for i in permutation[:train_examples]]
        test_set = [files[i] for i in permutation[-test_examples:]]
        train_labels = labels[permutation[:train_examples]]
        test_labels = labels[permutation[-test_examples:]]

        train_folder = datasetfolderpath + '/train'
        test_folder = datasetfolderpath + '/test'

        if isdir(train_folder):
            shutil.rmtree(train_folder)
        if isdir(test_folder):
            shutil.rmtree(test_folder)

        makedirs(train_folder + '/positive/')
        makedirs(train_folder + '/negative/')
        makedirs(test_folder + '/positive/')
        makedirs(test_folder + '/negative/')

        for f,i in zip(train_set, train_labels):
            if i == 0:
                shutil.copy2(f, train_folder+'/positive/')
            else:
                shutil.copy2(f, train_folder+'/negative/')

        for f,i in zip(test_set, test_labels):
            if i == 0:
                shutil.copy2(f, test_folder+'/positive/')
            else:
                shutil.copy2(f, test_folder+'/negative/')

In [7]: datagen = ImageDataGenerator(
        rescale = 1./255,
        rotation_range = 5,
        zoom_range = 0.2,
        horizontal_flip = True
    )

    img_height = img_width = 100

    channels = 3
    train_generator = datagen.flow_from_directory(
        train_folder,
        color_mode = 'rgb',

```

```

        target_size = (img_height, img_width),
        batch_size = 1,
        class_mode = None
    )

    i = 0
    img_list = []

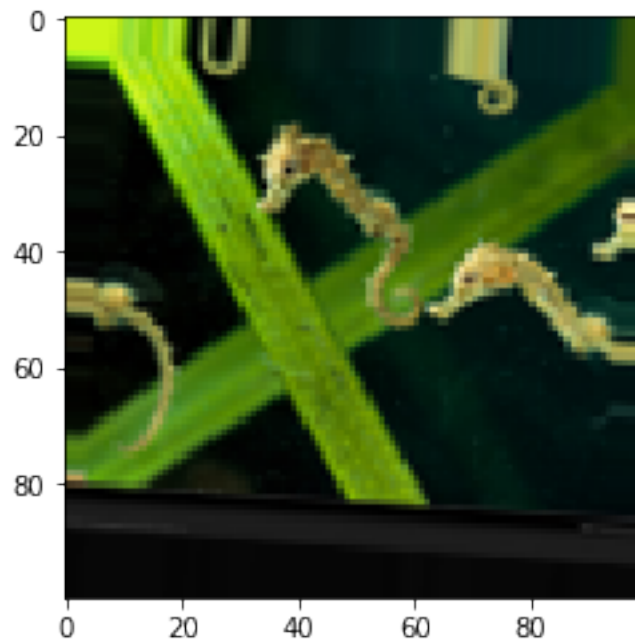
    for batch in train_generator:
        img_list.append(batch)
        i += 1


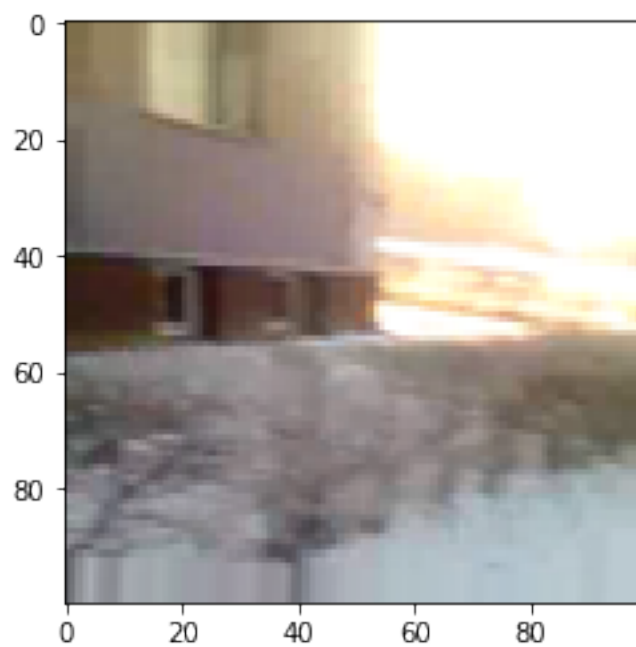
        if i > 5:
            break


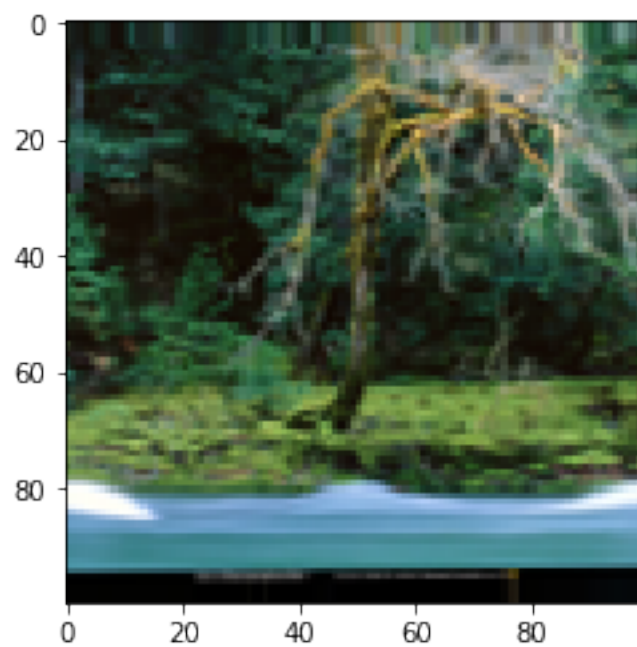
    for img in img_list:
        plt.imshow(np.squeeze(img))
        plt.show()


```

Found 10000 images belonging to 2 classes.









0.1.3 Logistic Regression

```
In [8]: batch_size = 1000
        train_generator = datagen.flow_from_directory(
            train_folder,
            color_mode = "rgb",
            target_size = (img_height, img_width),
            batch_size = batch_size,
            class_mode = 'binary'
        )
```

Found 10000 images belonging to 2 classes.

```
In [9]: x_train, y_train = next(train_generator)
        x_test, y_test = next(train_generator)
```

```
In [10]: from sklearn.linear_model import LogisticRegression

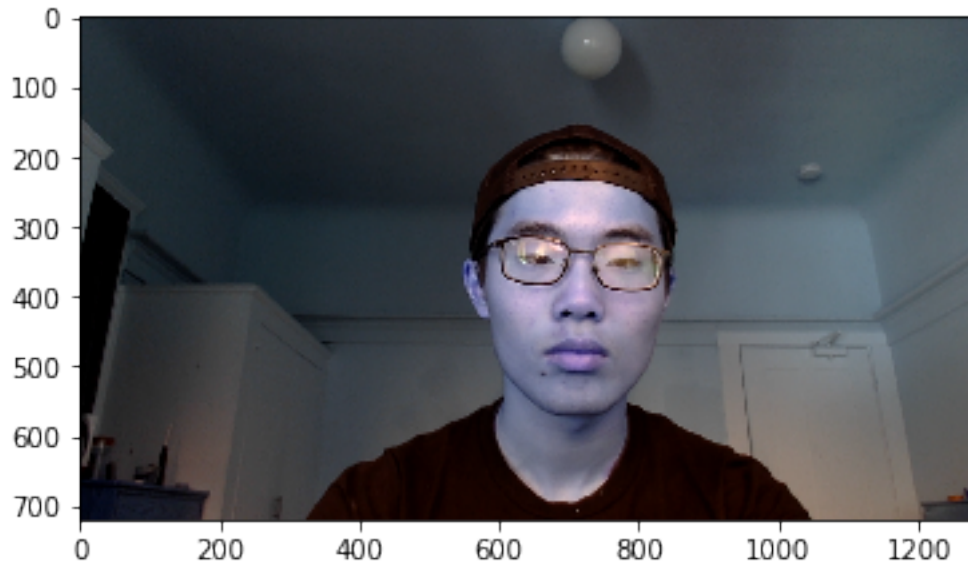
         logistic = LogisticRegression()
         logistic.fit(x_train.reshape(batch_size,-1), y_train)
```

```
Out[10]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                             penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                             verbose=0, warm_start=False)
```

```
In [11]: y_pred = logistic.predict(x_test.reshape(len(x_test), -1))
         y_pred[:10]

Out[11]: array([0., 0., 0., 0., 1., 0., 1., 1., 0., 0.], dtype=float32)

In [12]: plt.imshow(get_image())
         plt.show()
```



```
In [13]: image = get_image()
         cv2.imwrite('test.png', image)

Out[13]: True

In [14]: from PIL import Image

         img = Image.open('test.png')
         input_shape = (3,3)

         img_resized = img.resize(input_shape, Image.ANTIALIAS)
         img = np.asarray(img.resize((img_height, img_width), Image.ANTIALIAS))/255

In [15]: p = logistic.predict([img.reshape(-1)])
         if p[0] == 0.00:
             print("This is a knife")
         else:
             print("This is not a knife")

This is a knife

In [16]: np.count_nonzero(y_pred == y_test) / len(y_test)

Out[16]: 0.708
```

0.1.4 Convolutional Neural Networks (CNN)

```
In [17]: CNN_model = Sequential()
```

```
CNN_model.add(Conv2D(8, kernel_size=(3,3),padding='same',input_shape = (img_width, img_height, 3)))
CNN_model.add(Activation('relu'))
CNN_model.add(MaxPool2D(pool_size=(3,3)))

CNN_model.add(Conv2D(16, kernel_size=(3,3),padding='same'))
CNN_model.add(BatchNormalization())
CNN_model.add(Activation('relu'))
CNN_model.add(MaxPool2D(pool_size=(2,2)))

CNN_model.add(Conv2D(32, kernel_size=(3,3),padding='same'))
CNN_model.add(BatchNormalization())
CNN_model.add(Activation('relu'))
CNN_model.add(MaxPool2D(pool_size=(2,2)))

CNN_model.add(Conv2D(32, kernel_size=(3,3),padding='same'))
CNN_model.add(BatchNormalization())
CNN_model.add(Activation('relu'))
CNN_model.add(MaxPool2D(pool_size=(2,2)))

CNN_model.add(Flatten())
CNN_model.add(Dense(1, activation='sigmoid'))
CNN_model.compile(optimizer='adadelta',loss='binary_crossentropy',metrics=['accuracy'])
```

```
In [18]: CNN_model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 100, 100, 8)	224
activation_1 (Activation)	(None, 100, 100, 8)	0
max_pooling2d_1 (MaxPooling2D)	(None, 33, 33, 8)	0
conv2d_2 (Conv2D)	(None, 33, 33, 16)	1168
batch_normalization_1 (Batch Normalization)	(None, 33, 33, 16)	64
activation_2 (Activation)	(None, 33, 33, 16)	0
max_pooling2d_2 (MaxPooling2D)	(None, 16, 16, 16)	0
conv2d_3 (Conv2D)	(None, 16, 16, 32)	4640
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 32)	128


```

-----
activation_3 (Activation)      (None, 16, 16, 32)          0
-----
max_pooling2d_3 (MaxPooling2) (None, 8, 8, 32)            0
-----
conv2d_4 (Conv2D)              (None, 8, 8, 32)           9248
-----
batch_normalization_3 (Batch (None, 8, 8, 32)           128
-----
activation_4 (Activation)      (None, 8, 8, 32)            0
-----
max_pooling2d_4 (MaxPooling2) (None, 4, 4, 32)            0
-----
flatten_1 (Flatten)            (None, 512)                  0
-----
dense_1 (Dense)                 (None, 1)                    513
=====
Total params: 16,113
Trainable params: 15,953
Non-trainable params: 160
-----

```

```


In [19]: batch_size = 128
         train_generator = datagen.flow_from_directory(
             train_folder,
             color_mode = "rgb",
             target_size=(img_height, img_width),
             batch_size=batch_size,
             class_mode='binary')
         history = CNN_model.fit_generator(train_generator, train_examples//batch_size, epochs=10)

Found 10000 images belonging to 2 classes.
Epoch 1/10
78/78 [=====] - 144s 2s/step - loss: 0.4953 - acc: 0.7682
Epoch 2/10
78/78 [=====] - 128s 2s/step - loss: 0.4041 - acc: 0.8122
Epoch 3/10
78/78 [=====] - 150s 2s/step - loss: 0.3695 - acc: 0.8333
Epoch 4/10
78/78 [=====] - 133s 2s/step - loss: 0.3384 - acc: 0.8511
Epoch 5/10
78/78 [=====] - 141s 2s/step - loss: 0.3228 - acc: 0.8638
Epoch 6/10
78/78 [=====] - 99s 1s/step - loss: 0.3019 - acc: 0.8716
Epoch 7/10
78/78 [=====] - 79s 1s/step - loss: 0.2817 - acc: 0.8820
Epoch 8/10

```

```
78/78 [=====] - 78s 1s/step - loss: 0.2708 - acc: 0.8868
Epoch 9/10
78/78 [=====] - 86s 1s/step - loss: 0.2642 - acc: 0.8896
Epoch 10/10
78/78 [=====] - 85s 1s/step - loss: 0.2631 - acc: 0.8905
```

```
In [21]: plt.plot(history.history['acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
In [22]: batch_size = 1
test_generator = datagen.flow_from_directory(
    test_folder,
    color_mode = "rgb",
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='binary',
    shuffle=False)
y_pred = CNN_model.predict_generator(test_generator, test_examples//batch_size, worker
# model.predict_classes(test_x)
# np.count_nonzero(y_pred == test_y)/len(test_y)
```

Found 2899 images belonging to 2 classes.

```
In [23]: correct = 0
         true_correct = 0
         for i, f in enumerate(test_generator_filenames):
             if f.startswith('negative') and y_pred[i]<0.5:
                 correct +=1
             if f.startswith('positive') and y_pred[i]>=0.5:
                 correct +=1

         print('Correct predictions: '+str(correct/len(test_generator_filenames)))
```

Correct predictions: 0.8720248361503967

```
In [25]: batch_size = 6
         test_generator = datagen.flow_from_directory(
             test_folder,
             color_mode = "rgb",
             target_size=(img_height, img_width),
             batch_size=batch_size,
             class_mode='binary',
             shuffle=True)
         x_test, y_test = next(test_generator)






         p = CNN_model.predict(x_test)
         label_dict = {0: 'negative', 1: 'positive'}
```

Found 2899 images belonging to 2 classes.

```
In [26]: plt.figure(figsize=(35,35))

         for i in range(batch_size):
             plt.subplot(batch_size,2,2*i+1)
             plt.imshow(x_test[i])
             plt.title(label_dict[y_test[i]])

         plt.show()
```



```

In [48]: image = get_image()
         cv2.imwrite('test.png', image)

         img = Image.open('test.png')

         img = np.asarray(img.resize((img_height, img_width), Image.ANTIALIAS))/255


         r = CNN_model.predict(np.array([img]))

In [49]: plt.imshow(Image.open('test.png'))
         if r[0][0] < 0.5:
             print("This is not a knife")
             print("Confidence: ", 1 - r[0][0])

         else:
             print("This is a knife")
             print("Confidence: ", r[0][0])

```

This is a knife
Confidence: 0.8787047



0.1.5 Transfer Learning (VGG-16)

```

In [71]: from keras import applications

```

```

datagen = ImageDataGenerator(rescale=1.0/255)
VGG_model = applications.VGG16(include_top=False, input_shape=(img_width, img_height,

```

```

In [184]: batch_size = 128
          generator = datagen.flow_from_directory(
              train_folder,
              target_size=(img_width, img_height),
              batch_size=batch_size,
              class_mode=None,
              shuffle=False)
          #bottleneck_features_train = VGG_model.predict_generator(generator, train_examples//b
          #pickle.dump(bottleneck_features_train, open('bottleneck_features_train.npy', 'wb'))
          #bottleneck_features_train.shape

```

Found 10000 images belonging to 2 classes.

```

In [185]: batch_size = 128
          valid_generator = datagen.flow_from_directory(
              test_folder,
              target_size=(img_width, img_height),
              batch_size=batch_size,
              class_mode=None,
              shuffle=False)
          #bottleneck_features_valid = VGG_model.predict_generator(generator, test_examples//b
          #with open('bottleneck_features_valid.npy', 'wb') as f:
          #    pickle.dump(bottleneck_features_valid, f)
          #bottleneck_features_valid.shape

```

Found 2899 images belonging to 2 classes.

```

In [186]: with open('bottleneck_features_train.npy','rb') as f:
          bottleneck_features_train = pickle.load(f)

          VGG_model = Sequential()
          VGG_model.add(Flatten(input_shape = bottleneck_features_train.shape[1:]))
          VGG_model.add(Dense(256, activation='relu'))
          VGG_model.add(Dropout(0.5))
          VGG_model.add(Dense(1, activation='sigmoid'))

          VGG_model.compile(optimizer='rmsprop',loss='binary_crossentropy',metrics=['accuracy'])
          VGG_model.summary()

```

Layer (type)	Output Shape	Param #
flatten_5 (Flatten)	(None, 4608)	0

dense_5 (Dense)	(None, 256)	1179904

dropout_1 (Dropout)	(None, 256)	0

dense_6 (Dense)	(None, 1)	257
=====		
Total params: 1,180,161		
Trainable params: 1,180,161		
Non-trainable params: 0		

```
In [187]: batch_size = 128
          generator = datagen.flow_from_directory(
              train_folder,
              target_size=(img_width, img_height),
              batch_size=batch_size,
              class_mode=None,
              shuffle=False)


          labels = np.array([0 if f.startswith('negative') else 1 for f in generator.filesnames])
          VGG_history = VGG_model.fit(bottleneck_features_train, labels, epochs=10, batch_size=batch_size)
```

Found 10000 images belonging to 2 classes.

```
Epoch 1/10
9984/9984 [=====] - 3s 346us/step - loss: 0.4904 - acc: 0.8258
Epoch 2/10
9984/9984 [=====] - 3s 289us/step - loss: 0.2528 - acc: 0.8905
Epoch 3/10
9984/9984 [=====] - 3s 287us/step - loss: 0.2085 - acc: 0.9140
Epoch 4/10
9984/9984 [=====] - 3s 272us/step - loss: 0.1723 - acc: 0.9252
Epoch 5/10
9984/9984 [=====] - 3s 317us/step - loss: 0.1607 - acc: 0.9344
Epoch 6/10
9984/9984 [=====] - 3s 269us/step - loss: 0.1422 - acc: 0.9425
Epoch 7/10
9984/9984 [=====] - 3s 261us/step - loss: 0.1197 - acc: 0.9561
Epoch 8/10
9984/9984 [=====] - 3s 271us/step - loss: 0.1104 - acc: 0.9569
Epoch 9/10
9984/9984 [=====] - 3s 273us/step - loss: 0.0939 - acc: 0.9637
Epoch 10/10
9984/9984 [=====] - 4s 388us/step - loss: 0.0902 - acc: 0.9668
```

```
In [188]: plt.plot(VGG_history.history['acc'])
          plt.title('model accuracy')
```

```
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
In [189]: with open('bottleneck_features_valid.npy','rb') as f:
          bottleneck_features_valid = pickle.load(f)

          valid_labels = np.array([0 if f.startswith('negative') else 1 for f in valid_generator])
          y_valid_pred = VGG_model.predict_classes(bottleneck_features_valid)
          accuracy = np.count_nonzero(valid_labels == y_valid_pred.ravel())/len(valid_labels)

          print('\nThe accuracy is: '+str(accuracy))
```

The accuracy is: 0.7574573863636364

```
In [190]: vgg16 = applications.VGG16(include_top=False, weights='imagenet', input_shape=(img_w, img_h, 3))
          combinedModel = Model(inputs= vgg16.input, outputs= VGG_model(vgg16.output))

          for layer in combinedModel.layers[:-3]:
              layer.trainable = False
```

```
In [191]: combinedModel.summary()
```


Layer (type)	Output Shape	Param #
input_3 (InputLayer)	(None, 100, 100, 3)	0
block1_conv1 (Conv2D)	(None, 100, 100, 64)	1792
block1_conv2 (Conv2D)	(None, 100, 100, 64)	36928
block1_pool (MaxPooling2D)	(None, 50, 50, 64)	0
block2_conv1 (Conv2D)	(None, 50, 50, 128)	73856
block2_conv2 (Conv2D)	(None, 50, 50, 128)	147584
block2_pool (MaxPooling2D)	(None, 25, 25, 128)	0
block3_conv1 (Conv2D)	(None, 25, 25, 256)	295168
block3_conv2 (Conv2D)	(None, 25, 25, 256)	590080
block3_conv3 (Conv2D)	(None, 25, 25, 256)	590080
block3_pool (MaxPooling2D)	(None, 12, 12, 256)	0
block4_conv1 (Conv2D)	(None, 12, 12, 512)	1180160
block4_conv2 (Conv2D)	(None, 12, 12, 512)	2359808
block4_conv3 (Conv2D)	(None, 12, 12, 512)	2359808
block4_pool (MaxPooling2D)	(None, 6, 6, 512)	0
block5_conv1 (Conv2D)	(None, 6, 6, 512)	2359808
block5_conv2 (Conv2D)	(None, 6, 6, 512)	2359808
block5_conv3 (Conv2D)	(None, 6, 6, 512)	2359808
block5_pool (MaxPooling2D)	(None, 3, 3, 512)	0
sequential_5 (Sequential)	(None, 1)	1180161
Total params: 15,894,849		
Trainable params: 3,539,969		
Non-trainable params: 12,354,880		

0.1.6 Fine-tuning (This may take one to two hours)

```
In [101]: VGG_model.save_weights('fc_model.h5')
```

```
In [102]: combinedModel.compile(loss='binary_crossentropy',  
                                optimizer = optimizers.RMSprop(lr=1e-4, decay=0.9), # optimizers.SGD(lr=1e-4),  
                                metrics=['accuracy'])
```

```
# prepare data augmentation configuration
```

```
train_datagen = ImageDataGenerator(  
    rescale=1. / 255,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True)
```

```
test_datagen = ImageDataGenerator(rescale=1. / 255)
```

```
train_generator = train_datagen.flow_from_directory(  
    train_folder,  
    target_size=(img_height, img_width),  
    batch_size=batch_size,  
    class_mode='binary')
```

```
validation_generator = test_datagen.flow_from_directory(  
    test_folder,  
    target_size=(img_height, img_width),  
    batch_size=batch_size,  
    class_mode='binary')
```

Found 10000 images belonging to 2 classes.

Found 2899 images belonging to 2 classes.

```
In [103]: # fine-tune the model
```

```
combinedModel.fit_generator(  
    train_generator,  
    steps_per_epoch=train_examples//batch_size,  
    epochs=5,  
    validation_data=validation_generator,  
    validation_steps=test_examples//batch_size) # len(validation_generator.filenames)
```

WARNING:tensorflow:Variable *= will be deprecated. Use variable.assign_mul if you want assignment

Epoch 1/5

78/78 [=====] - 2334s 30s/step - loss: 1.2813 - acc: 0.4110 - val_loss: 1.2813

Epoch 2/5

78/78 [=====] - 4183s 54s/step - loss: 0.8882 - acc: 0.5181 - val_loss: 0.8882

Epoch 3/5

78/78 [=====] - 1626s 21s/step - loss: 0.8141 - acc: 0.5469 - val_loss: 0.8141

Epoch 4/5

```
78/78 [=====] - 1796s 23s/step - loss: 0.7804 - acc: 0.5663 - val_loss: 0.7804
Epoch 5/5
78/78 [=====] - 1873s 24s/step - loss: 0.7629 - acc: 0.5818 - val_loss: 0.7629
```

```
Out[103]: <keras.callbacks.History at 0x1a30c5cb00>
```

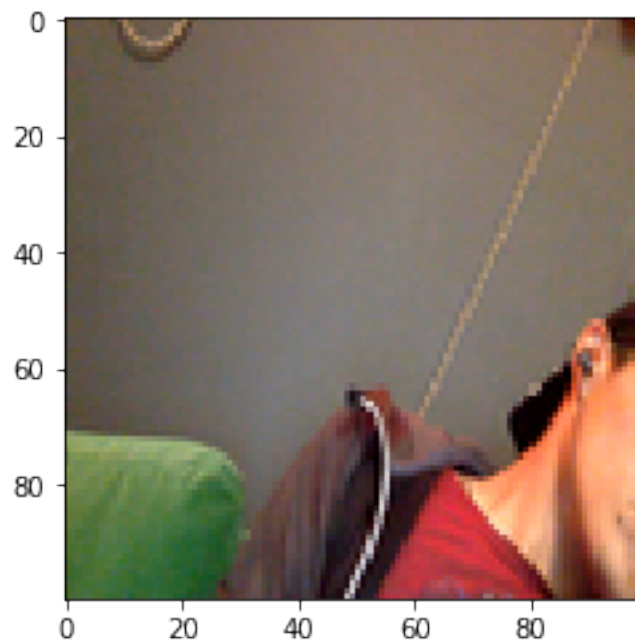
```
In [199]: image = get_image()
cv2.imwrite('test.png', image)
```

```
Out[199]: True
```

```
In [200]: img = Image.open('test.png')
input_shape = (3,3)

img_resized = img.resize(input_shape, Image.ANTIALIAS)

In [201]: img = Image.open('test.png')
img = np.asarray(img.resize((img_height, img_width), Image.ANTIALIAS))/255
plt.imshow(img)
plt.show()
p = combinedModel.predict(np.array([img]))
if p[0][0] < 0.5:
    print("This is not a knife")
    print("Confidence: ", 1-p[0][0])
else:
    print("This is a knife")
    print("Confidence: ", p[0][0])
```



This is a knife
Confidence: 0.97202975

```
In [5]: path = "/anaconda3/lib/python3.6/site-packages/cv2/data/"
        cam = cv2.VideoCapture(0)
        detector=cv2.CascadeClassifier(path + 'haarcascade_frontalface_default.xml')

In [ ]: while(True):
        ret, img = cam.read()
        gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
        faces = detector.detectMultiScale(gray, 1.3, 5)
        for (x,y,w,h) in faces:
            cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

        cv2.imshow('frame',img)
        if cv2.waitKey(1) & 0xFF == ord('q'):
            break
```